

Kronecker Products and GPT2

Ayoub @ Uni Passau

April 8, 2024

Abstract

LLMs are cool, everywhere, but are pretty big.

Contents

1	Introduction	3
2	Methodology	3
2.1	Kronecker Decomposition	3
2.2	The 95M model	3
2.3	Different decomposition schemes:	4
2.4	Initialization	5
2.4.1	Van Loan Method	6
2.4.2	Pruning based Initialization	6
3	Experiments	6
3.1	Training setup	6
3.2	Results	6
3.3	The 81M class:	7
3.4	The 95M class:	7
3.5	General trends:	7
4	Discussion	7
4.1	Parameter count:	8
4.2	Only 3% of data	8

1 Introduction

A lot of work has been done on the compression of LLMs using factorization methods, including Kronecker Products. Most notably for encoder based, and rarely applied to decoder based methods ([1], [2], [3]).

We study how far we can push the compression ratio, and record then impact of down-scaling on the quality of the model. We also investigate adding multiple Kronecker factors.

Our contributions (so far, delete this) could be summarized as follows:

- We use weight tying, i.e., the embedding and softmax matrices are identical. This makes our model, the only "effective" 81M model compared to (X,Y).
- We have a systematic compression scheme, i.e., compress all layers the same way. We don't see any reasons to why we would only compress odd layers (besides to "force" the number of parameters to be under certain threshold).
- We try different compression schemes (67M, 81M, and 95M)
 - We propose a new simple initialization for the 95M model.
 - We use multiple factors for the VL based inits.
- We don't use distillation (empirical evidence shows no benefits compared to only Vanilla supervised learning).

2 Methodology

2.1 Kronecker Decomposition

In this section, we study different Kronecker decomposition setups, and the percentage of compression it would lead to. So far we only decompose the weights `c_fc.weight` and `c_proj.weight` (each has 2.3M in the original GPT2-small architecture.). Each transformer layer (there are 12 in total) has two of these weights. They count to 56.6M in total (45% of GPT2 124M). Hence any significant reduction to these matrices would lead to a remarkable compression ratio of the whole model. We choose not to compress the other weights, namely, attention weights and embedding matrix for various reasons that we will expose later on.

2.2 The 95M model

The most basic strategy is to divide one of the dimensions of each W by 2, this would lead to a 95M model. The parameter $p_1 = \text{c_fc.weight}$ (resp. $p_2 = \text{c_proj.weight}$) has a shape of (3072, 768) (resp. (768, 3072)). We first try the following decomposition: $p_1 = \underbrace{W_{11}}_{(3072,384)} \otimes \underbrace{W_{12}}_{(1,2)}$ and $p_2 =$

$$\underbrace{W_{21}}_{(384,3072)} \otimes \underbrace{W_{22}}_{(2,1)}$$

This decomposition would lead to to reduction of 28M (50%). The new network would have approx 95M. Our goal is to eventually reach the

82M mark, similar to DistilGPT2, and other Factorized models (inserts other refs here).

2.3 Different decomposition schemes:

It is reasonable to aim for a decomposition that guarantees the maximum rank we can get. Since the Rank of Kronecker products is multiplicative, i.e., $\text{rank}(A \otimes B) = \text{rank}(A) \cdot \text{rank}(B)$ ¹, we can easily compute the rank of each possible decomposition. In our case, we have $W \in R^{(m,n)}$ where $m = 3072$ and $n = 768$. Hence, for each layer of GPT2, we aim to find the “optimal” $A \in R^{(m_1, n_1)}$, and $B \in R^{(m_2, n_2)}$, i.e.,:

$$W \approx A \otimes B, \quad m = m_1 m_2, n = n_1 n_2$$

W.l.o.g, for each decomposition (A, B) the maximum rank we can reach is $\min(m_1, n_1) \cdot \min(m_2, n_2)$. And each of the reduced decompositions would have exactly $m_1 n_1 + m_2 n_2$ parameters. Hence, theoretically, the maximum rank we can get is 768 of a $(3072, 768)$ matrix. The following table summarizes some possible combinations, alongside the reduction it would lead to per layer, and the total number of parameters in GPT2, for only those decompositions of maximal attainable rank. We are particularly interested in 3 class of models, the 67M, the 81M and the 96M. (Need to add this) Furthermore, we add multiple factors to the models labeled with **MF** (second table / a few decompositions are missing, check this out. e.g., 3072, 384).

¹Link to proof: [https://math.stackexchange.com/questions/707091/elementary-proof-for-rank\(A \otimes B\) = \text{rank}\(A\) \cdot \text{rank}\(B\)](https://math.stackexchange.com/questions/707091/elementary-proof-for-rank(A \otimes B) = \text{rank}(A) \cdot \text{rank}(B))

Name	Dimension	params	Model size
67M	(64, 32)	3200	67,893,504
	(64, 48)	3840	67,908,864
	(96, 32)	3840	67,908,864
	(64, 64)	4672	67,928,832
	(128, 32)	4672	67,928,832
	(96, 48)	5120	67,939,584
	(96, 64)	6528	67,973,376
	(128, 48)	6528	67,973,376
	(128, 64)	8480	68,020,224
	(96, 96)	9472	68,044,032
	(192, 48)	9472	68,044,032
	(128, 96)	12480	68,116,224
	(192, 64)	12480	68,116,224
	(128, 128)	16528	68,213,376
MF1	(256, 64)	16528	68,213,376

MF2	(1024, 256)	262153	74,108,376
	(768, 384)	294920	74,894,784
	(1024, 384)	393222	77,254,032
81M	(768, 768)	589828	81,972,576
	(1536, 384)	589828	81,972,576
	(1024, 768)	786435	86,691,144
96M	(1536, 768)	1179650	96,128,304
GPT2	(3072, 768)	2359297	124,439,832

Table 1: Different compression schemes

Name	Dimension	params	1 factor	2 factors	3 factors
MF1	(256, 64)	16528	68,2M	68,6	69M
MF2	(1024, 256)	262153	74M	80M	86M

Table 2: Adding multiple Kronecker Factors

2.4 Initialization

Since we inherit a GPT2 checkpoint that was trained for multiple epochs on the Open Web Text (OWT) (cite here), we want to initialize our weights in a way that leverages the old pre-training as much as possible. This is of course obvious for the parameters that are common between **GPT2** and **KronyPT** (i.e., we match). But more tricky for the weights that are decomposed into Kronecker Factors. In our work, we try two different approaches, Van Loan decomposition (cite here), and a Pruning based method exclusively for the 95M model.

# Params	Model	Datasets		
		wikitext-103	wikitext-2	Lambada
95M	Krony-PT1	41.80	35.50	61.34
95M	Krony-PT1	41.81	36.02	59.95

Table 3: Comparison of different models on various datasets.

2.4.1 Van Loan Method

XX (maybe not even write this section)

2.4.2 Pruning based Initialization

We propose a new initialization strategy by inducing sparsity in the first factor of the Kronecker Product, and prune it by half. This is equivalent to picking the second factor as $\begin{bmatrix} 1 \\ 0 \end{bmatrix}$. Now, we illustrate how this procedure works with a random matrix.

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \\ a_{41} & a_{42} & a_{43} & a_{44} & a_{45} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \\ a_{61} & a_{62} & a_{63} & a_{64} & a_{65} \end{bmatrix} \xrightarrow{\text{pruning}} \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ 0 & 0 & 0 & 0 & 0 \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \\ 0 & 0 & 0 & 0 & 0 \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \\ 0 & 0 & 0 & 0 & 0 \\ a_{61} & a_{62} & a_{63} & a_{64} & a_{65} \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \\ a_{61} & a_{62} & a_{63} & a_{64} & a_{65} \end{bmatrix} \otimes \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

3 Experiments

3.1 Training setup

For pre-training, we follow general industry standards, namely of Chinchilla [4]. We noticed more stable loss curves with higher **batch sizes**, typically 128 batch size (attained with a gradient accumulation of 4, and each forward pass seeing a batch of 32), and constant **learning rate** of 6×10^{-5} (better than a cosine scheduler with a warm-up). After pre-training on approximately one epoch, we pick the best checkpoint and train further on ≈ 500 k tokens per step (as suggested in the literature) which amounts in our set up (sequence length of 1024) to a batch size of 512. All experiments have been conducted using a single A100 80 GB, for one epoch of on Open Web Text [5], an open source reproduction of OpenAI’s WebText dataset.

3.2 Results

We refer to our models with the nomenclature **Krony-PT-XM{-Y}**, with X representing the number of the parameters, and optionally Y refer-

ring to separate checkpoints of the same model class. We namely focus on 2 class models, the **Krony-PT-81M** and **Krony-PT-95M**. We evaluate our models next-token prediction capabilities on 3 datasets, wikitext103, wikitext2 [6] and Lambada [7].

3.3 The 81M class:

Krony-PT-81M is our model class of 81M parameters, and the suffixes 1350 and 3950 refer to different checkpoints. Both models clearly outperform distilGPT2 ([8]) on the 3 datasets, with a noticeable better perplexity on Lambada [7].

# Params	Model	Datasets		
		wikitext-103	wikitext-2	Lambada
124M	GPT2	29.16	24.67	45.28
82M	DistilGPT2	44.53	36.48	76.00
81M	KronyPT-81M-1350	41.98	34.99	-
81M	KronyPT-81M-3950	-	-	64.92

Table 4: Perplexity results of Krony-PT and DistilGPT

# Params	Model	Datasets		
		wikitext-103	wikitext-2	Lambada
81M	KronyPT-81M-1350	41.98	34.99	-
81M	KronyPT-81M-3950	-	-	64.92
119M	TQCompressedGPT2	40.28	32.25	64.72
119M	KnGPT-2	40.97	32.81	67.62

Table 5: Perplexity of Krony-PT against other Kronecker based models.

3.4 The 95M class:

3.5 Scale test

How far can we push down the size? tl;dr: 67M does not work!

3.6 General trends:

Things to add here:

- Models do not get better at Wiki, but do improve on Lambada.
- Distillation is not helpful
- Does pruning help? For training and for inference?
- does re-plugging the matrix work?

4 Discussion

We clarify in this section some design choices of the other papers:

4.1 Parameter count:

For some reason the other papers do not include the output matrix, and I quote for

4.2 Only 3% of data

both papers claim that they only used 3% of the training data to train their models. We argue that in this set up, limiting your training data to only a fraction does not imply a better compression/factorization method, for the simple and following reasons:

1. They inherit almost half the weights from the old GPT2 checkpoint that was trained for numerous epochs, and has seen the whole data.
2. They use Van Loan method, hence, even, when they don't match the original checkpoint, VL is still a smart init, some knowledge is definitely passed through the SVD.
3. You literally use the output matrix of a pre-trained GPT2.

Hence, we don't know "how much knowledge" has been passed through the already trained parameters. A fair comparison would be to initialize **all the parameters** of the new compressed model with random values, and not rely on any of the other pre-trained ones. Which is clearly not the case here.

References

- [1] M. Tahaei, E. Charlaix, V. Nia, A. Ghodsi, and M. Rezagholizadeh, “Kroneckerbert: Significant compression of pre-trained language models through kronecker decomposition and knowledge distillation,” in *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pp. 2116–2127, 2022.
- [2] A. Edalati, M. Tahaei, A. Rashid, V. P. Nia, J. J. Clark, and M. Rezagholizadeh, “Kronecker decomposition for gpt compression,” *arXiv preprint arXiv:2110.08152*, 2021.
- [3] V. Abronin, A. Naumov, D. Mazur, D. Bystrov, K. Tsarova, A. Melnikov, I. Oseledets, S. Dolgov, R. Brasher, and M. Perelshtein, “Tq-compressor: improving tensor decomposition methods in neural networks via permutations,” *arXiv preprint arXiv:2401.16367*, 2024.
- [4] J. Hoffmann, S. Borgeaud, A. Mensch, E. Buchatskaya, T. Cai, E. Rutherford, D. de Las Casas, L. A. Hendricks, J. Welbl, A. Clark, *et al.*, “Training compute-optimal large language models. arxiv,” *arXiv preprint arXiv:2203.15556*, 2022.
- [5] A. Gokaslan and V. Cohen, “Openwebtext corpus.” <http://Skylion007.github.io/OpenWebTextCorpus>, 2019.
- [6] S. Merity, C. Xiong, J. Bradbury, and R. Socher, “Pointer sentinel mixture models,” *arXiv preprint arXiv:1609.07843*, 2016.
- [7] D. Paperno, G. Kruszewski, A. Lazaridou, Q. N. Pham, R. Bernardi, S. Pezzelle, M. Baroni, G. Boleda, and R. Fernández, “The lambda dataset: Word prediction requiring a broad discourse context,” *arXiv preprint arXiv:1606.06031*, 2016.
- [8] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter,” in *NeurIPS EMC²Workshop*, 2019.